

Reflection: Building My Website with Generative AI

Nischal Sadashivaiah (中文名：拉吉 / Raj) · M1561031 · Generative AI HW1 · 494 words

Live site: <https://nischal-sadashivaiah.vercel.app> · Source: <https://github.com/NischalRaj16/nischal-portfolio-2>

For this assignment I used Claude as a design and coding collaborator to build and publish my personal academic website. I gave it the full homework handout and asked it to help me build a site that a professor could understand in about two minutes.

What I asked the AI to do. I asked it to turn the handout into concrete requirements, write the HTML, CSS and JavaScript, create the visuals, test the layout, and deploy the site to Vercel. Before writing any code, it asked me four questions about hosting, which research to feature, visual style, and which contact details to make public. This made me decide on my audience first. I first chose a clean academic style for a professor audience, then pushed it toward a colourful, modern tech design.

What it did well. In one pass the AI produced a responsive single-page layout with semantic HTML, a mobile menu, and accessibility details such as a skip link and alt text. The pipeline diagram it drew for my intrusion detection paper explains my research faster than a paragraph can. It was honest: it did not invent an accuracy figure, dataset or graduation year, and it asked me to confirm the Chinese name it suggested.

Weaknesses and errors I found. The first draft looked finished, but it had three real bugs that only appeared when I tested it at phone width. A LinkedIn entry that I had marked as hidden was still visible, because a CSS display rule overrode the hidden attribute. This would have been a broken link. The university name broke in the middle of a word, and the wide research diagram shrank until its text was unreadable. Reading the code alone would not catch these.

What I changed. I fixed the hidden-element rule and limited word wrapping to the long email address. I also added a vertical version of the diagram for small screens. On content, I chose to feature only my published paper instead of also listing ongoing lab work, I rejected the AI-generated Chinese name and used 拉吉 (Raj), the name I go by, replaced a generic research summary with my real results, and I kept my phone number and address off the public page. After deployment, I checked that every asset and external link returned HTTP 200. I also noticed that the preview URL required a Vercel login, so I submit the public production URL instead.

What I learned about human-AI collaboration. The AI is good at producing a solid first version quickly, but "looks right" is not the same as "works right." My job moved from typing code to setting requirements, checking facts, and testing on real screen sizes.

The most important decisions, like who the audience is, what is true about my research, and what personal information should be public, had to be mine. I now treat AI output like a draft from a capable teammate: useful, but only trusted after I have verified it.